

SPC : A tool to diagnose bottleneck  
component in a back end server  
software stack.

BPS-2 presentation by  
Jatin Lachhwani  
(Guide : Prof. Bhaskaran Raman)

# Contents

---

Introduction

SPC Overview

SPC Implementation

SPC Usage Testing

Conclusion & Future Work

# Introduction


# Introduction

## Background | **SAFE**

---

SAFE - Smart Authenticated  
Fast Exams

Take quiz in classroom using  
their own smartphones.




# Introduction

## Background | **SAFE** | Core Functionality

---

SAFE - Smart Authenticated  
Fast Exams

**Take quiz** in classroom using  
their own smartphones.



# Introduction

## Background | SAFE | Architecture

SAFE app

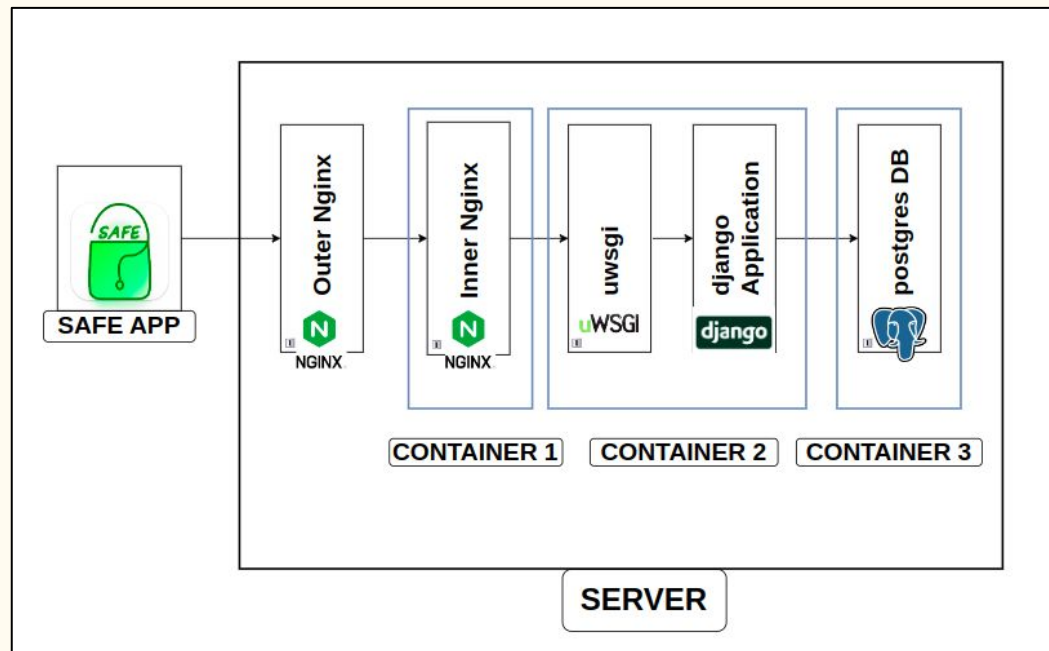
Outer nginx

Inner nginx

uWSGI

Django

Postgres



# Introduction

## Background | **SAFE** | Core Functionality

---

### Steps : **Taking a Quiz in SAFE**

0. Instructor starts the quiz



1. **Login**



2. Get **Course List**



3. Get **Quiz List**



4. Get **Quiz Info**



5. **Quiz Download**



6. **Quiz Authenticate**



7. **Quiz Submit**



# Introduction

## Background | Previous Work

---

To identify bottleneck APIs among them.

0. Instructor starts the quiz



1. **Login**



2. Get **Course List**



3. Get **Quiz List**



4. Get **Quiz Info**



5. **Quiz Download**



6. **Quiz Authenticate**



7. **Quiz Submit**





## Background | Previous Work

Tried to find bottleneck APIs using Locust.

API	Avg. Resp. time(ms)
1. Login	4279
2. Get Course List	4378
3. Get Quiz List	4278
4. Get Quiz Info	9407
5. Quiz Download	7766
6. Quiz Authenticate	4278
7. Quiz Submit	5704

Test Conditions

Num of users : 100

Spawn Rate : 10/s

# Introduction

## Background | Previous Work

Tried to find bottleneck APIs using Locust.

API	Avg. Resp. time(ms)
1. Login	4279
2. Get Course List	4378
3. Get Quiz List	4278
4. Get Quiz Info	9407
5. Quiz Download	7766
6. Quiz Authenticate	4278
7. Quiz Submit	5704

Too high

for these conditions

Test Conditions

Num of users : 100

Spawn Rate : 10/s

# Introduction

## Background | Previous Work

Tried to find bottleneck APIs using Locust.

API	Avg. Resp. time(ms)
1. Login	4279
2. Get Course List	4378
3. Get Quiz List	4278
4. Get Quiz Info	9407
5. Quiz Download	7766
6. Quiz Authenticate	4278
7. Quiz Submit	5704

Too high

for these conditions

Possibility of a system  
level bottleneck

Test Conditions

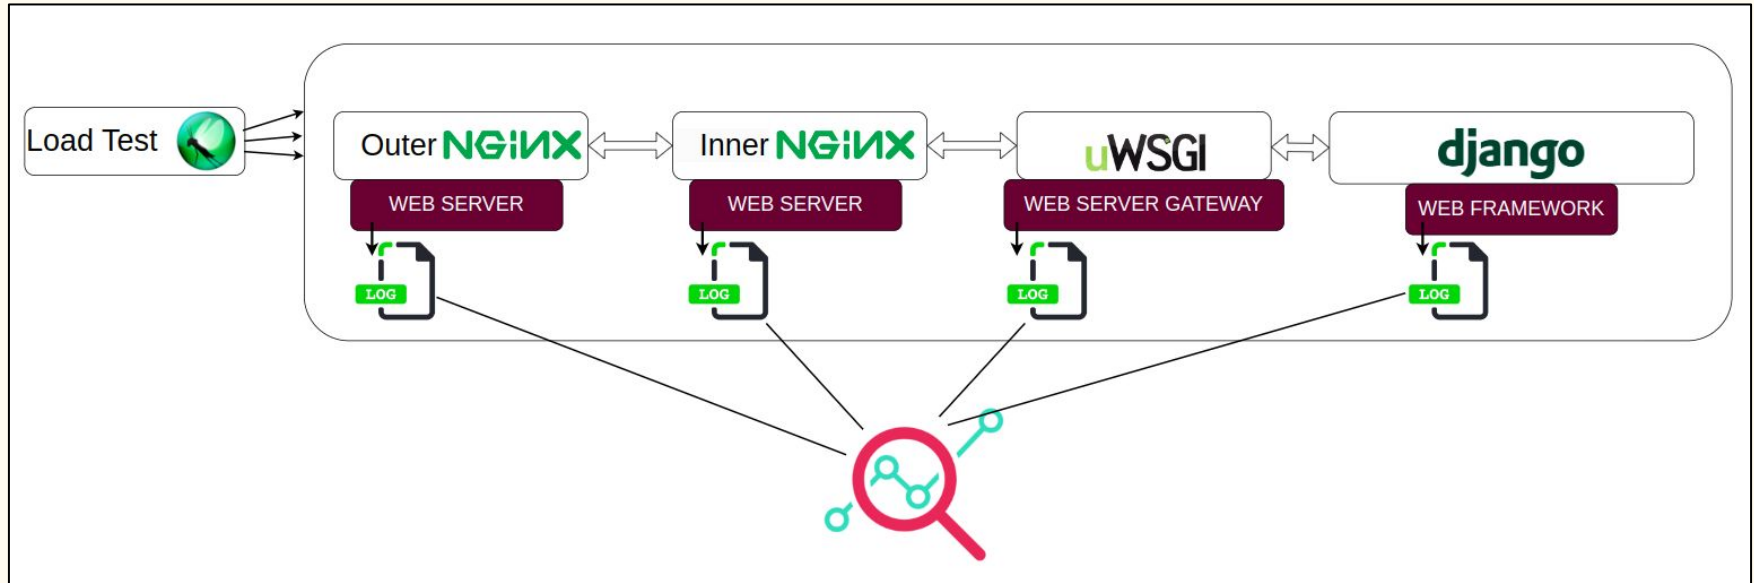
Num of users : 100

Spawn Rate : 10/s

# Introduction

## Background | Previous Work

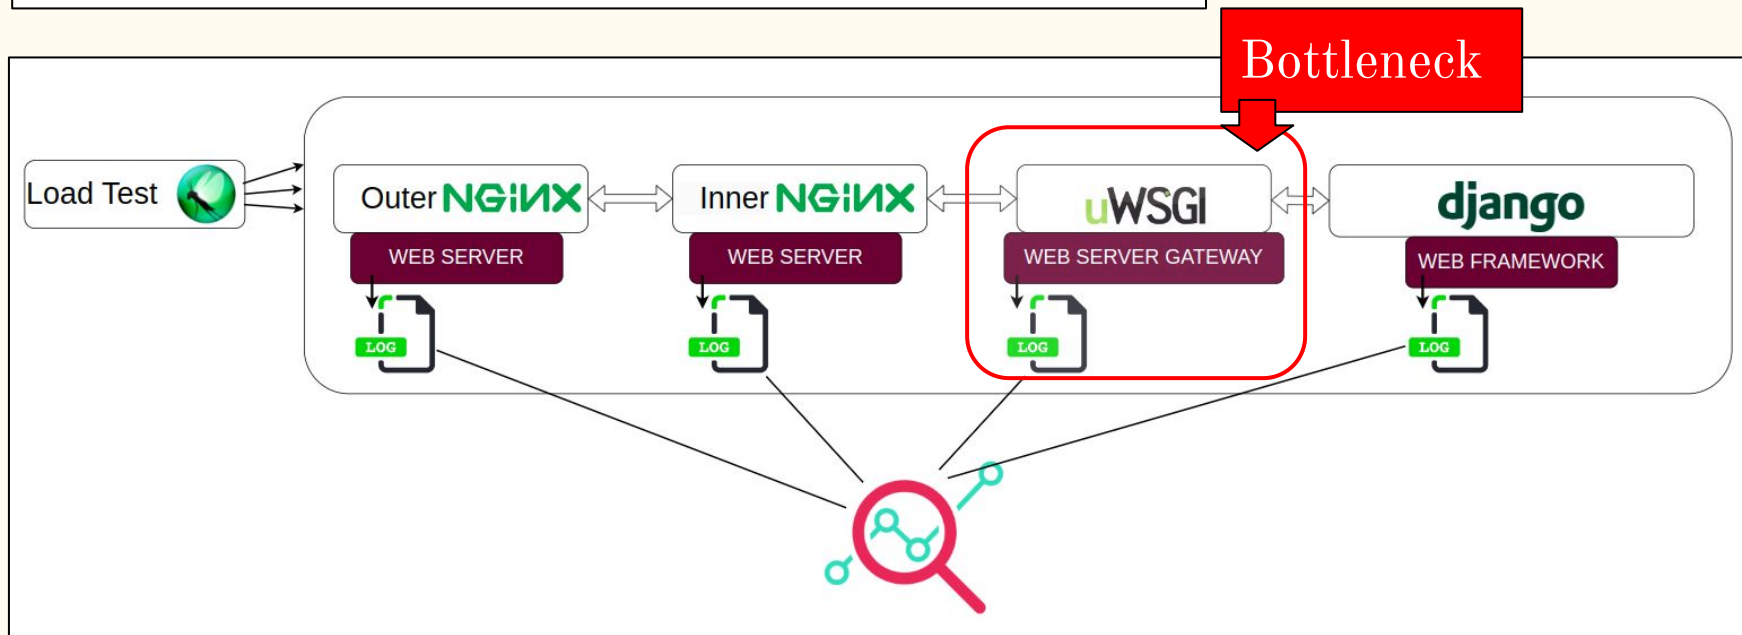
Analysed the logs of each component in SAFE



# Introduction

## Background | Previous Work

Analysed the logs of each component in SAFE



# Introduction

## Background | Previous Work

After Properly configuring **uWSGI** (bottleneck component)

API	Avg. Resp. time(ms)	Avg. Resp. time(ms) (uWSGI configured)
1. <b>Login</b>	4279	<b>1968</b>
2. Get <b>Course List</b>	4378	<b>438</b>
3. Get <b>Quiz List</b>	4278	<b>519</b>
4. <b>Get Quiz Info</b>	9407	<b>306</b>
5. <b>Quiz Download</b>	7766	<b>325</b>
6. <b>Quiz Authenticate</b>	4278	<b>405</b>
7. <b>Quiz Submit</b>	5704	<b>783</b>

Test Conditions

Num of users : 100

Spawn Rate : 10

# Introduction

## Background | Previous Work Conclusion

---

System level bottleneck component(s)

are hard to detect

can stay unnoticed for too long

have a significant impact on overall performance

## Problem Statement

---

**To create a tool that is**

able to **detect System level bottleneck component(s).**

**independent** of any **application specific APIs**

**easy to use and configure**

**visual feedback** for analysis of components.



# SPC Overview

## Assumptions

---

Possible to customize the log format for all components in the software stack.

Components are stacked linearly (like nginx-uwsgi-django)

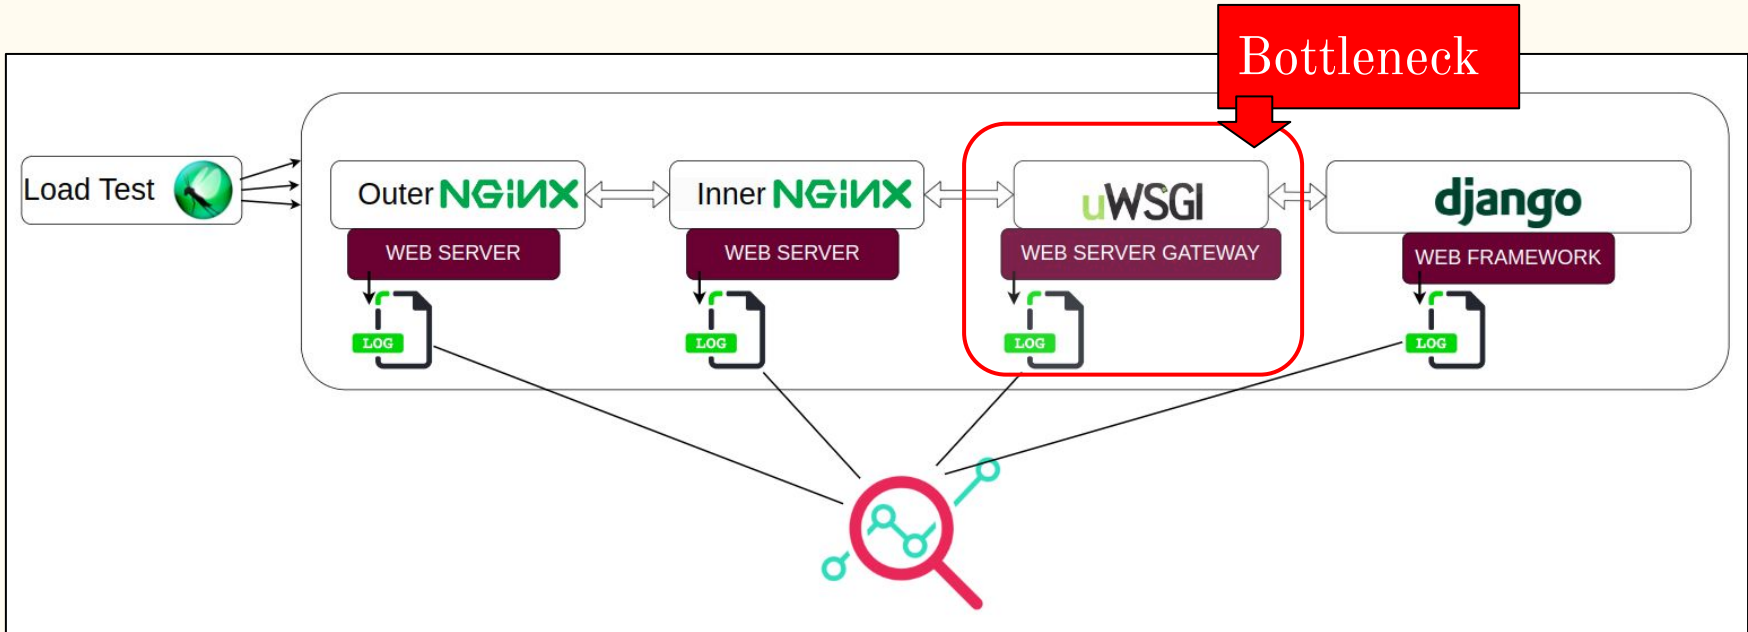
URL of the request is visible at all components.

Mainly designed for backend web application stacks.

# SPC Overview

## Test for SPC

**Good test for SPC :** It can detect the uWSGI bottleneck we detected in our previous work.



# SPC Overview

## SPC Design Brainstorm

---

How can we detect a **bottleneck** like **ill-configured uWSGI**?

while ensuring

**independent** of any **application specific APIs**

**easy to use and configure**

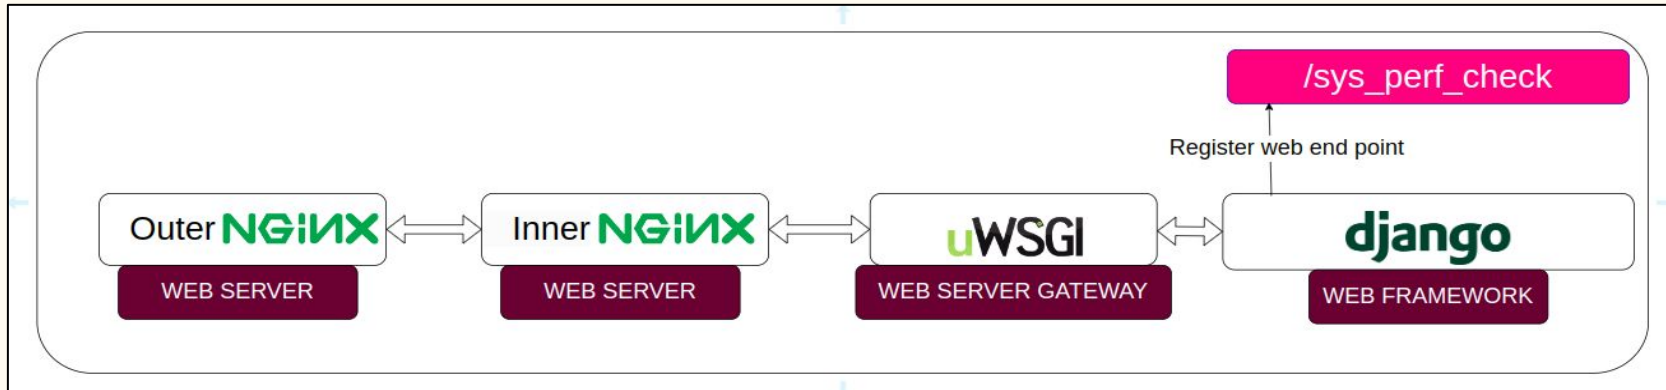
**visual feedback** for analysis of components.

# SPC Overview

## SPC Design Brainstorm | Idea 1

Add a specific API endpoint at application code independent of other APIs.

Say : sys\_perf\_check




# SPC Overview

## SPC Design Brainstorm | Idea 1

Add a specific API endpoint at application code independent of other APIs.

Say : sys\_perf\_check

Why?



# SPC Overview

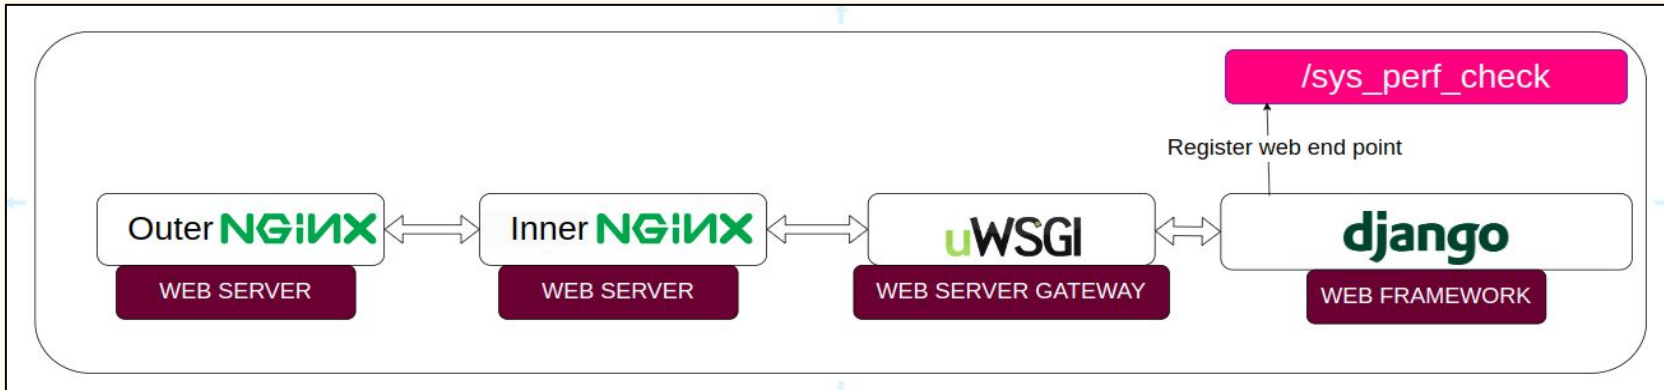
## SPC Design Brainstorm | Idea 1

Add a specific API endpoint at application code independent of other APIs.

Say : `sys_perf_check`

Why?

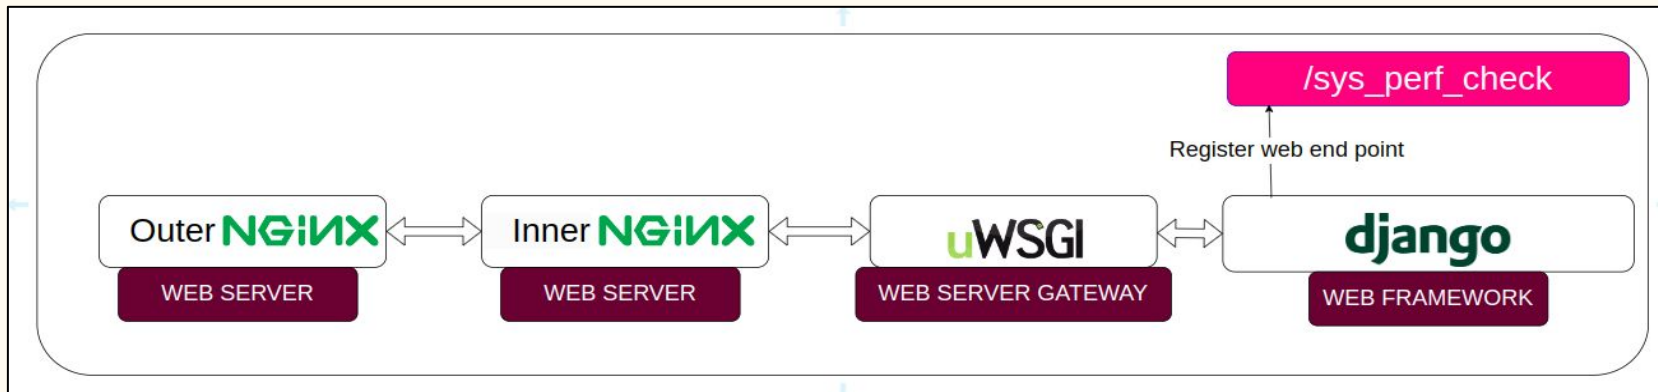
**independent of any application specific APIs**



# SPC Overview

## SPC Design Brainstorm | Idea 1

What will `/sys_perf_check` do?



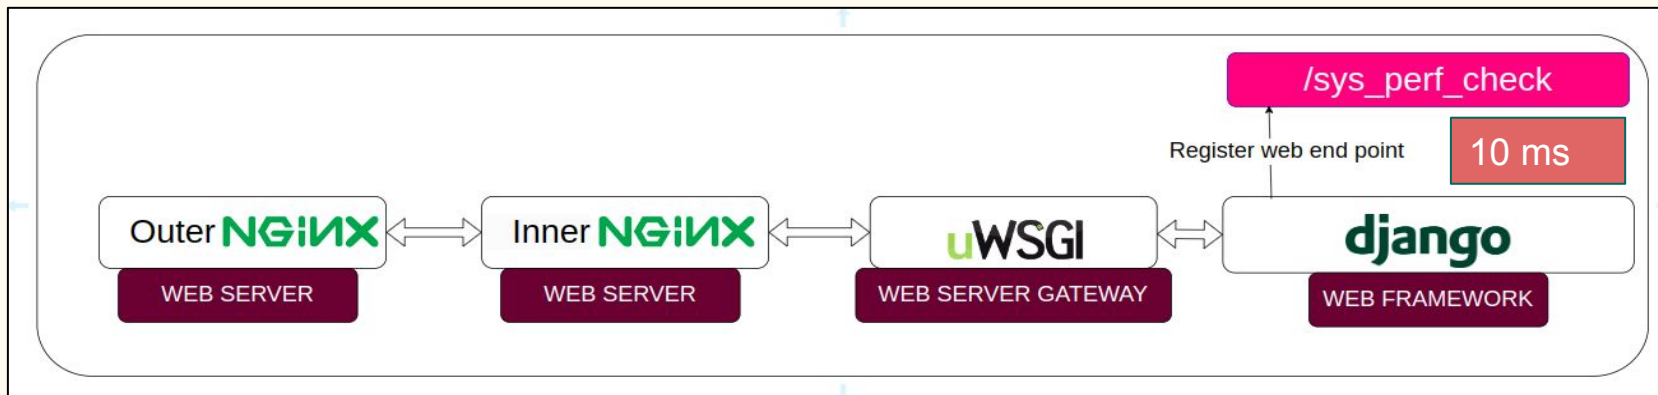


# SPC Overview

## SPC Design Brainstorm | Idea 1

What will `/sys_perf_check` do?

It will execute for a bounded time then exit. (say **10ms**)




# SPC Overview

## SPC Design Brainstorm | Idea 1

What will `/sys_perf_check` do?

It will execute for a bounded time then exit. (say **10ms**)

Why?



# SPC Overview

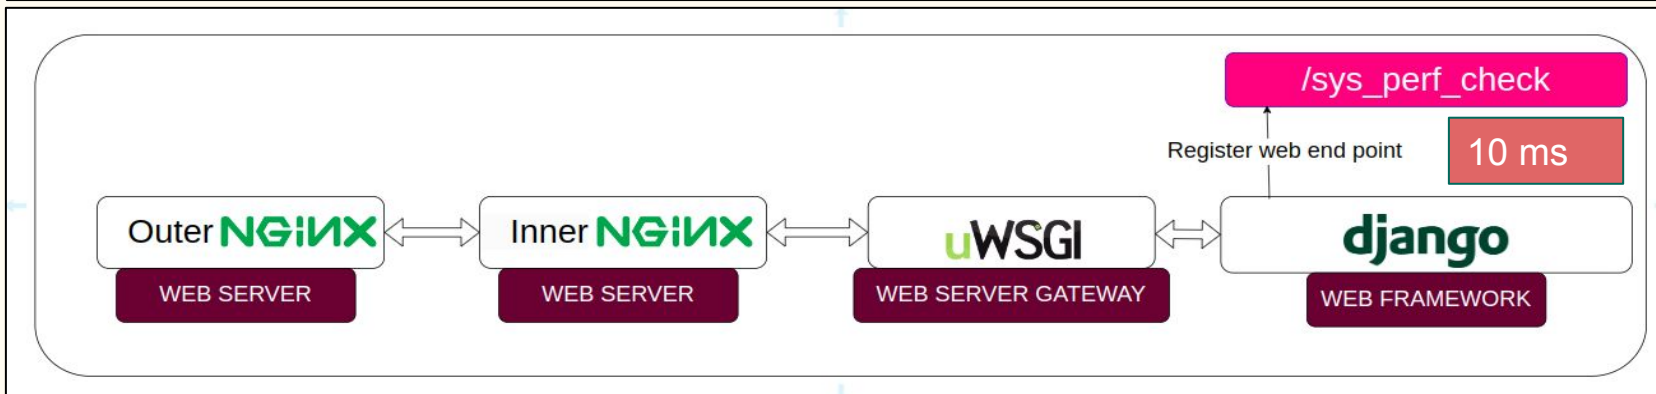
## SPC Design Brainstorm | Idea 1

What will `/sys_perf_check` do?

It will execute for a bounded time then exit. (say **10ms**)

Why?

We observed in our previous work, **django code** would **execute** for roughly **XY ms.** (empirical). (Adjusted based upon the system under test.)

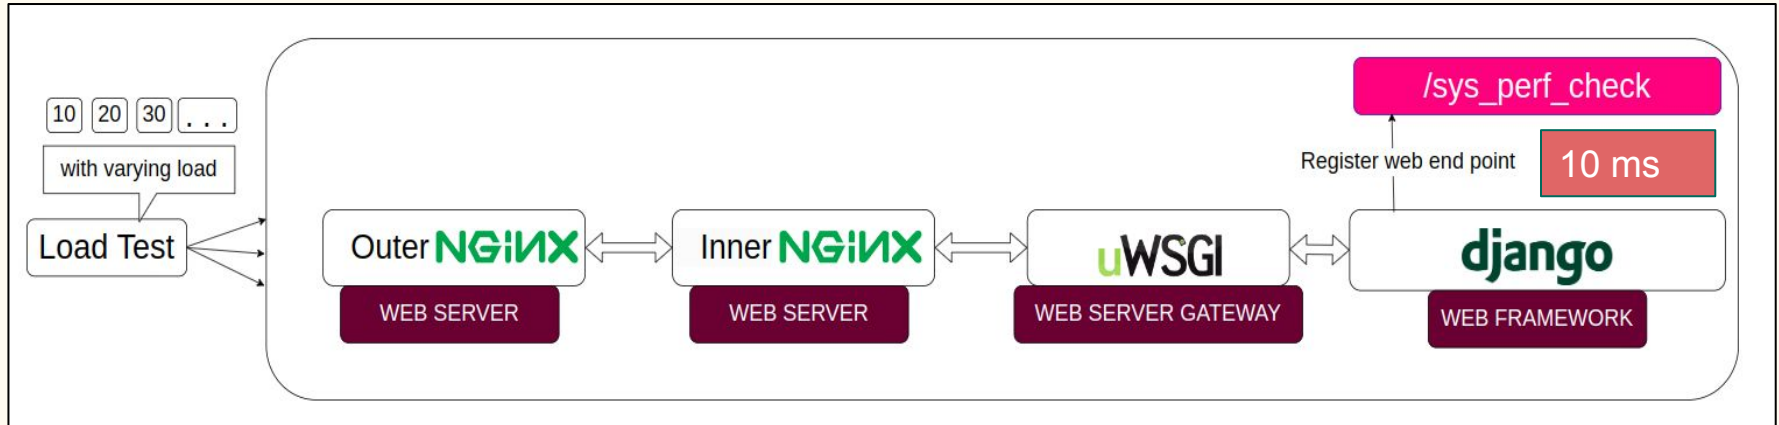


# SPC Overview

## SPC Design Brainstorm | Idea 1

Now we can do a **load test** with **varying number of users**.

We can **analyze the component logs** and **find the bottleneck**.



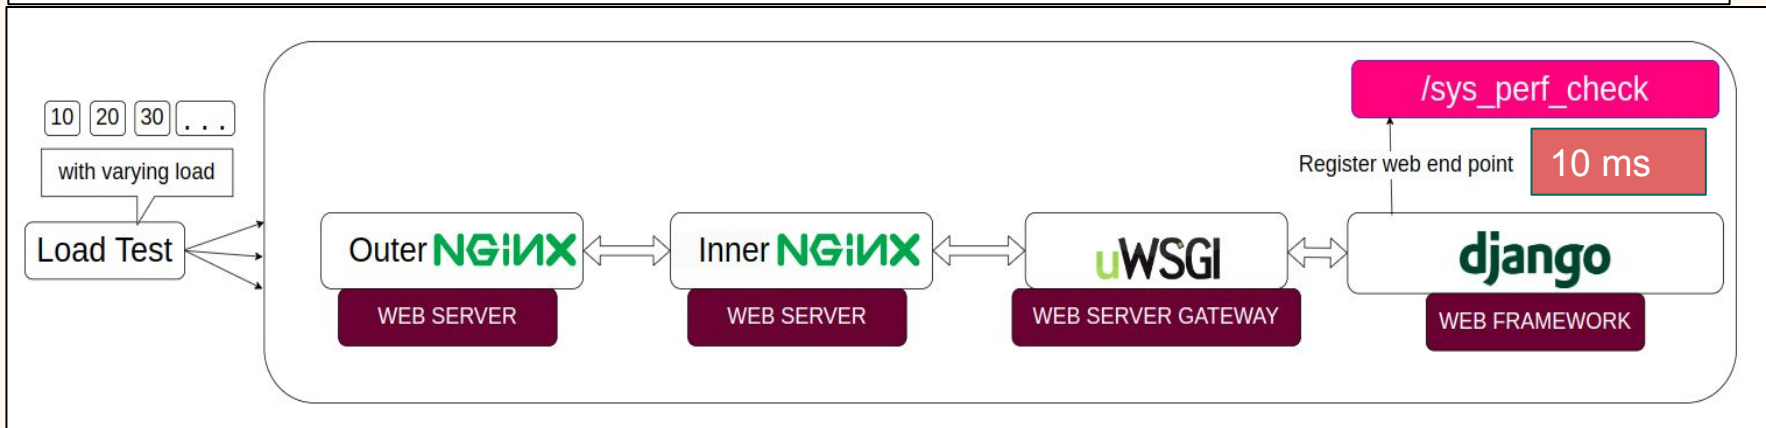
# SPC Overview

## SPC Design Brainstorm | Idea 2

Now we can do a **load test** with **varying number of users**.

We can **analyze the component logs** and **find the bottleneck**.

**Wait.** But **there is a catch.** How to identify the number of users for a request entry in a log?



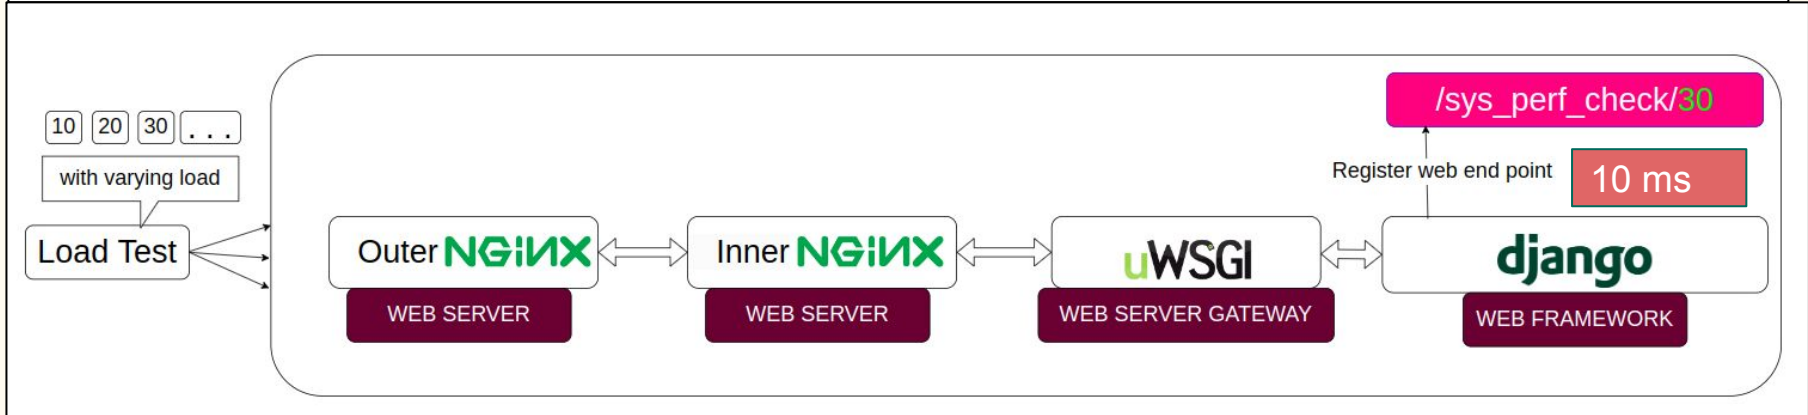
# SPC Overview

## SPC Design Brainstorm | Idea 2

Now we can do a **load test** with **varying number of users**.

We can **analyze the component logs** and **find the bottleneck**.

**Wait.** But **there is a catch.** How to identify the number of users for a request entry in a log? By **appending number of users as a part of request.**

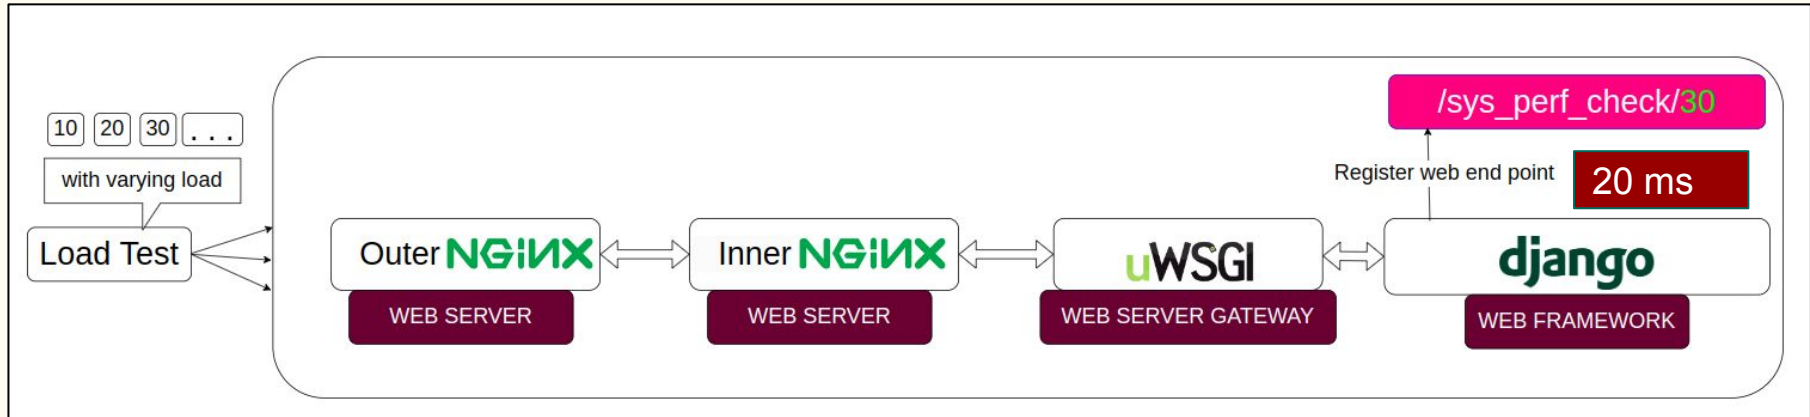


# SPC Overview

## SPC Design Brainstorm | Idea 3

Suppose we change something for a new test.(say endpoint time to 20 ms)

After the load test when we extract the logs. We might get logs from the previous test.

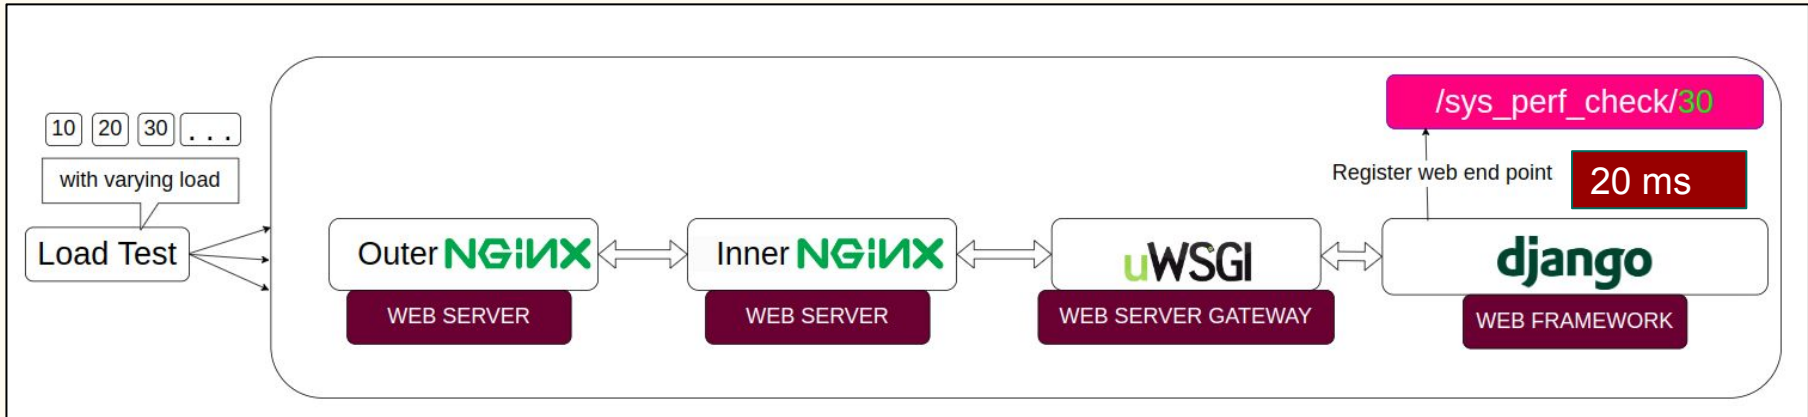


# SPC Overview

## SPC Design Brainstorm | Idea 3

Suppose we change something for a new test.(say endpoint time to 20 ms)

After the load test when we extract the logs. We might get logs from the previous test. What now?





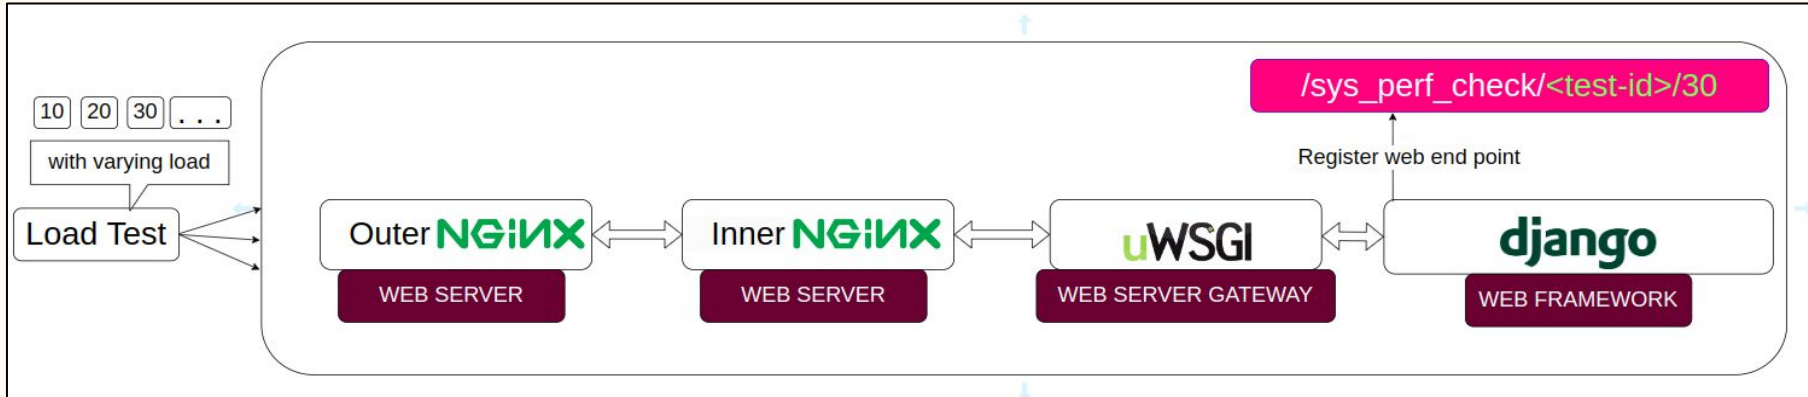
# SPC Overview

## SPC Design Brainstorm | Idea 3

Suppose we change something for a new test.(say endpoint time to 20 ms)

After the load test when we extract the logs. We might get logs from the previous test. **What now?** We need a way to distinguish tests uniquely

We will add a unique test id in the request for each varying load test.

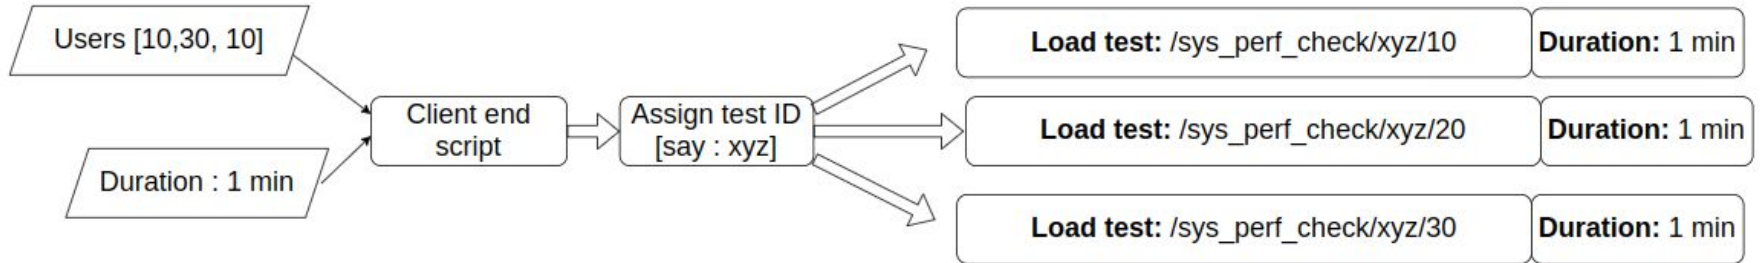


# SPC Overview

## SPC Design Brainstorm | Idea Aggregation

1. Register endpoint `sys_perf_check`
2. Script that generates load for varying number of users
3. Assigns Id to each instance of test.

These things will be done by a script : **client end script**



# SPC Overview

## SPC Design Brainstorm | Client end script

---

client end script

Why?

# SPC Overview

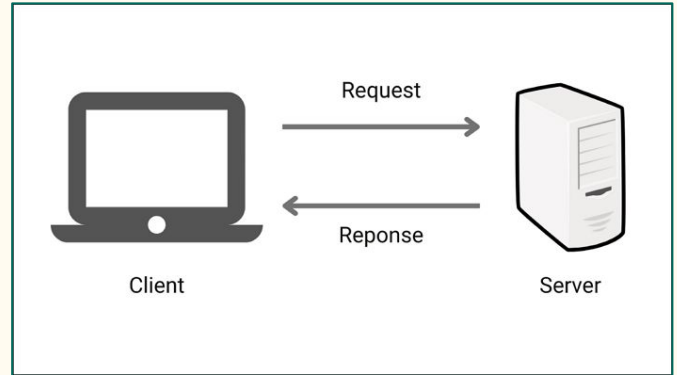
## SPC Design Brainstorm | Client end script

client end script

Why?

Generally **load generation for testing** happens on a **separate machine**.

To keep the **server machine** same for all tests.



## SPC Design Brainstorm | Idea 4

---

Once the client end script executes

We need something to

**extract the logs**

**transfer them to client end script for analysis**



## SPC Design Brainstorm | Idea 4

---

Once the client end script executes

We need something to

**extract the logs**

**transfer them to client end script for analysis**

**Server end script**



## SPC Design Brainstorm | Idea 4

---

Once the client end script executes

We need something to

**extract the logs**

**transfer them to client end script for analysis**

**Server end script**

SPC: **easy to use** and **configure**



## SPC Design Brainstorm | Idea 5

---

Once logs are received : Analyze them




# SPC Overview

## SPC Design Brainstorm | Idea 5

Once logs are received : Analyze them

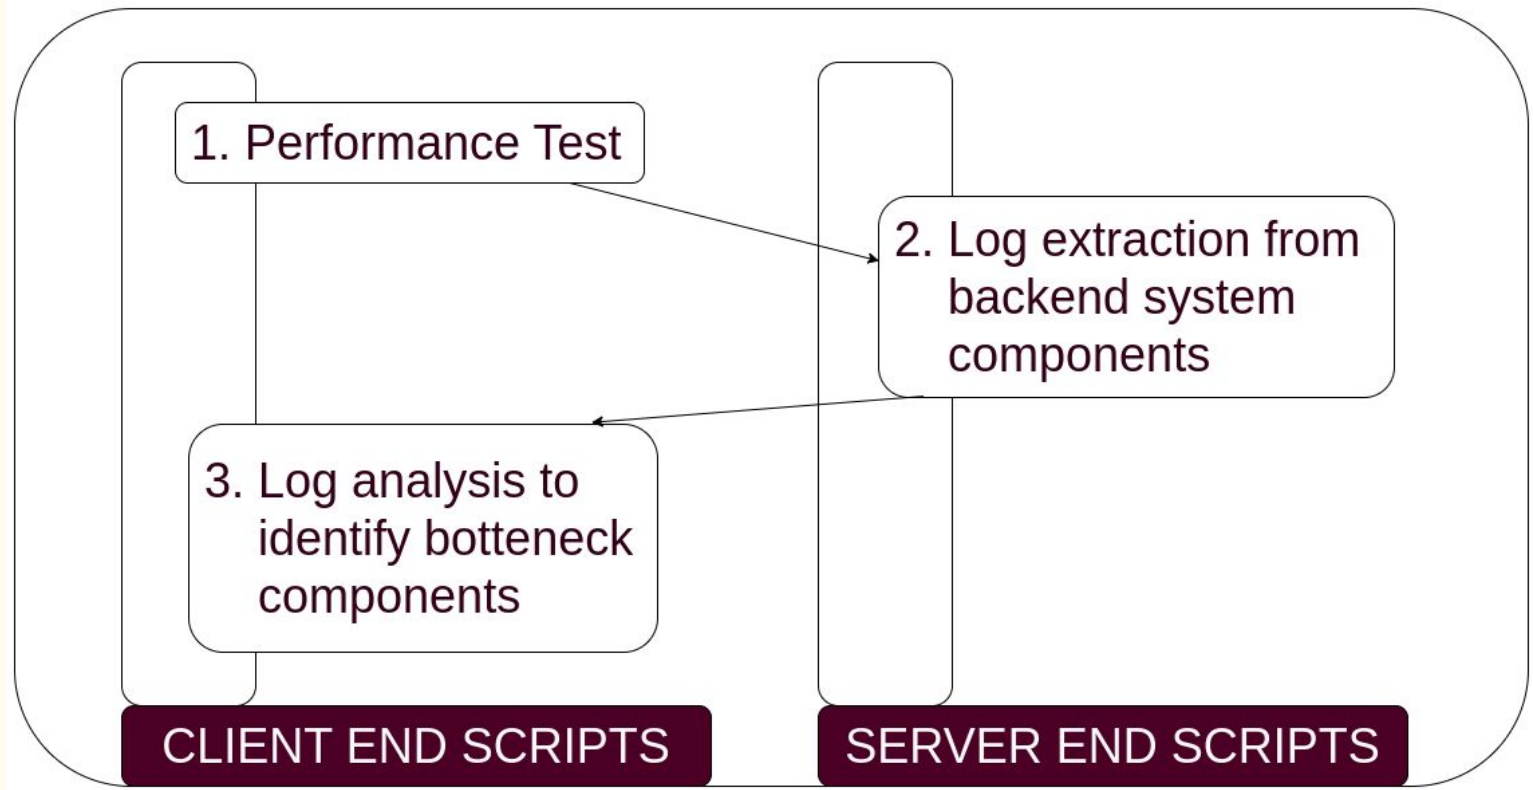
Display graphical results for analysis.

visual feedback for analysis of components.



## SPC Design

---

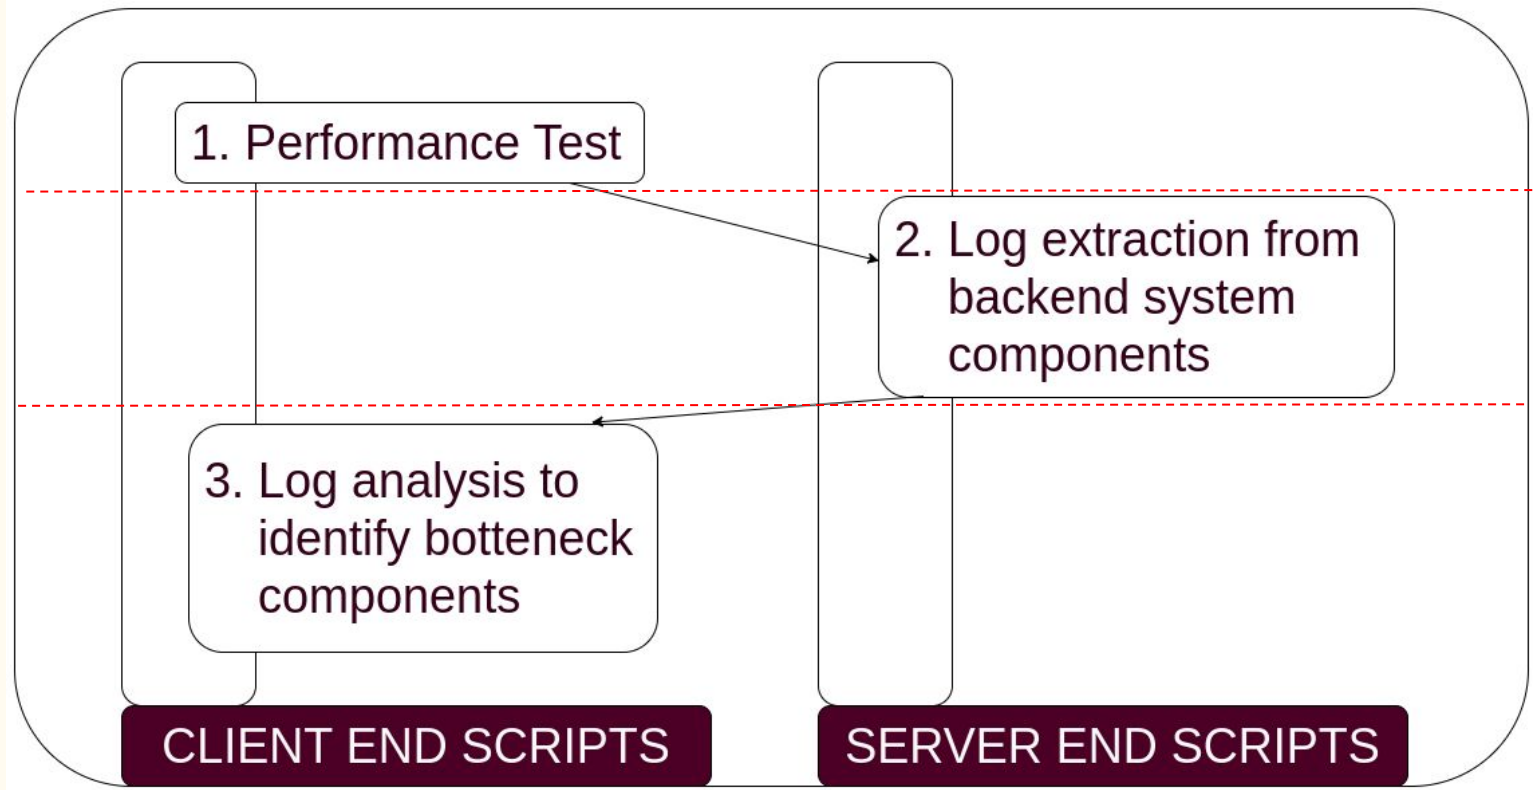


# SPC Implementation

# SPC Implementation

## 3 Phases

---



# SPC Implementation

## Phase 1 - Performance Test | Client End Script

---

### 1. Client end script input

```
λ > ./client_end_script.py -h thesis/locust_fi
usage: ./client_end_script.py [-h] -l LOWER_BOUND_USERS -u UPPER_BOUND_USERS -s STEP_SIZE
                             [-t RUN_TIME]
```

To figure out system level bottleneck component for REST API based services

optional arguments:

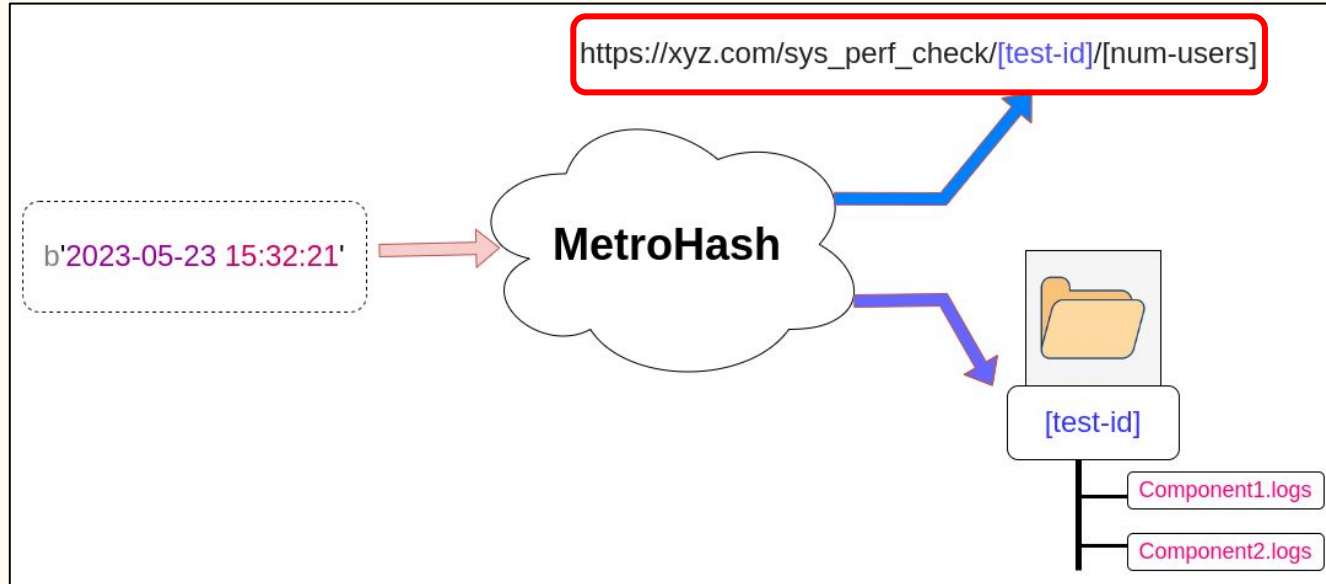
-h, --help	show this help message and exit
-l LOWER_BOUND_USERS	Specify the lower bound of the number of users.
-u UPPER_BOUND_USERS	Specify the upper bound of the number of users.
-s STEP_SIZE	Specify the step size for incrementing the number of users.
-t RUN_TIME	Specify the runtime for each user number being tested

# SPC Implementation

## Phase 1 - Performance Test | Client End Script

---

2. Generating unique id for this instance of script.

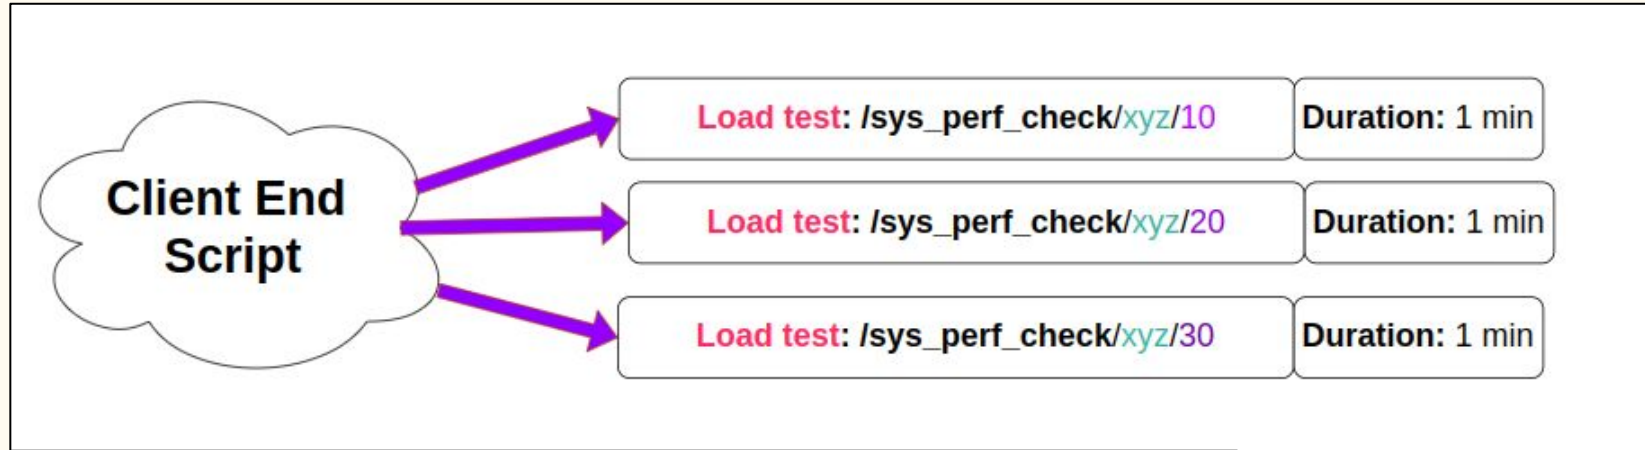


# SPC Implementation

## Phase 1 - Performance Test | Client End Script

---

### 3. Performance Test (uses locust)



say : l: 10, u:30, s:10, d:1 min | test-id : xyz

# SPC Implementation

## Transition Phase 1 to Phase 2

---

After Phase 1, client end script will **send the test-id** to server end script.

Client End Script

Send : Test-id

Server End script

```
sequenceDiagram
    participant Client as Client End Script
    participant Server as Server End script
    Client->>Server: Send : Test-id
```

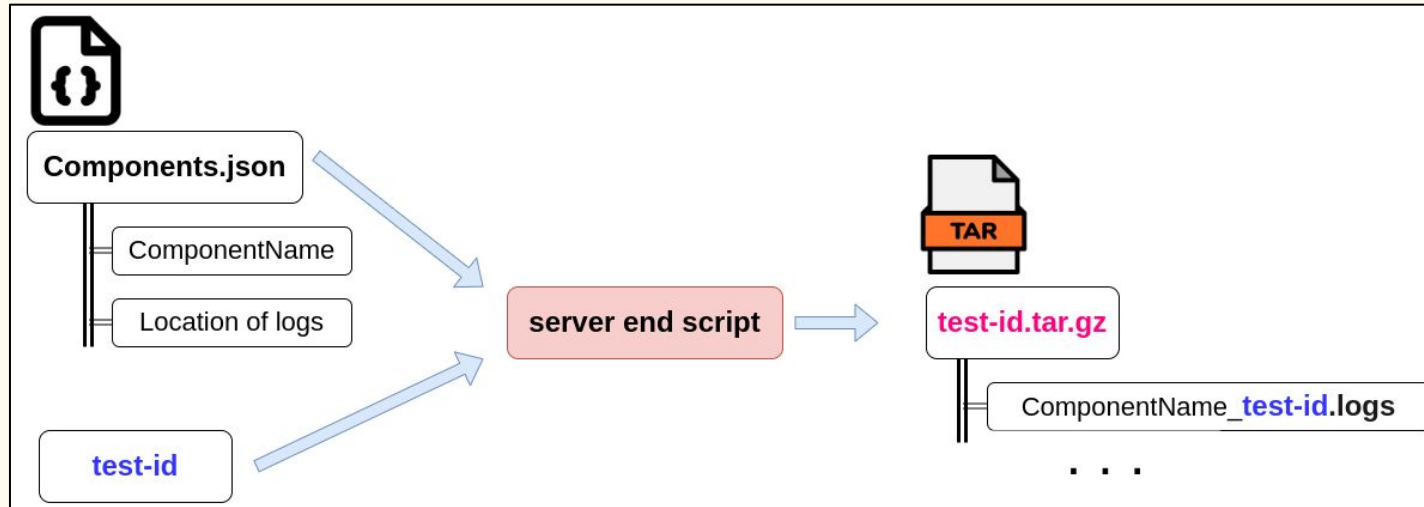


# SPC Implementation

## Phase 2 - Log Extraction | Server End Script

### 4. Log Extraction from Component logs.

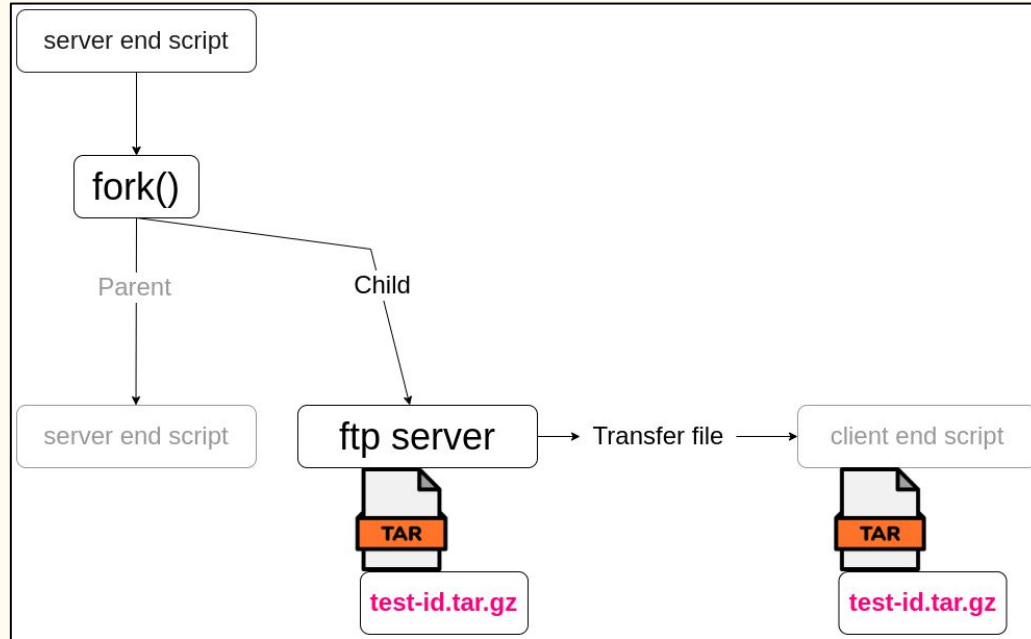
It requires a **components.json** file to describe the location of component logs.



# SPC Implementation

## Phase 2 - Log Extraction | Server End Script

### 5. Creating an ftp server to transfer logs

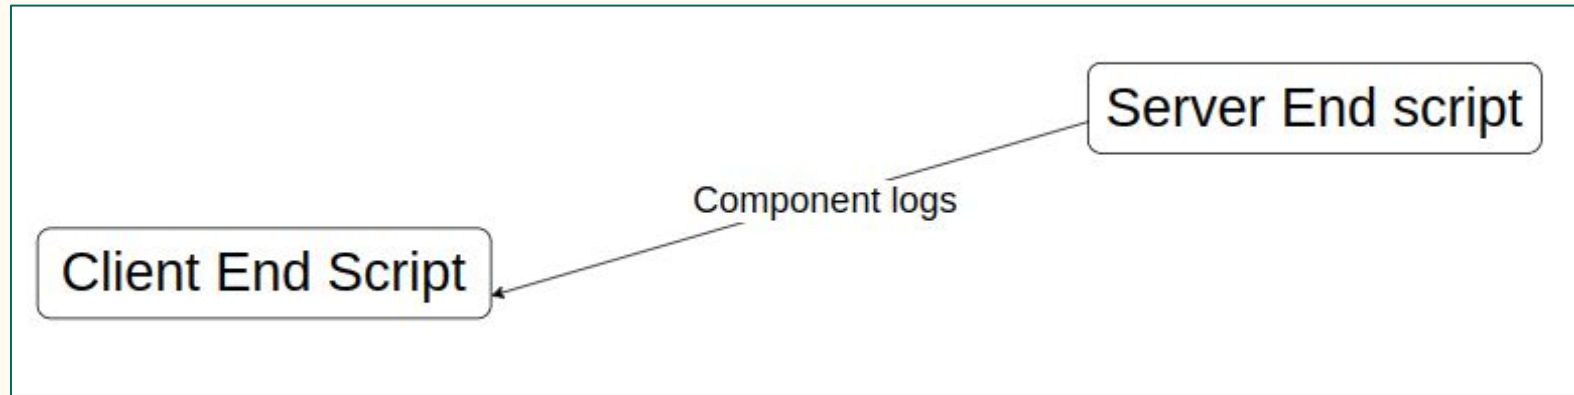


# SPC Implementation

## Transition Phase 2 to Phase 3

---

After Phase 2, Components logs are transferred to client end script.




# SPC Implementation

## Phase 3 - Graph for analysis| Client End Script

6. Component logs are used to create graph for analysis.

**Image and Table Presentation**



**Table containing Response Time(ms) vs number of users**

Numusers	outer-nginx_time(ms)	inner-nginx_time(ms)	uwsgi_time(ms)
10	25.41	24.69	24.56
20	25.97	25.29	25.11
30	26.61	25.86	25.58

This can be used to detect the bottleneck.

# SPC Usage Testing

# SPC Usage Testing

## Test Setup | Machine Specifications

---

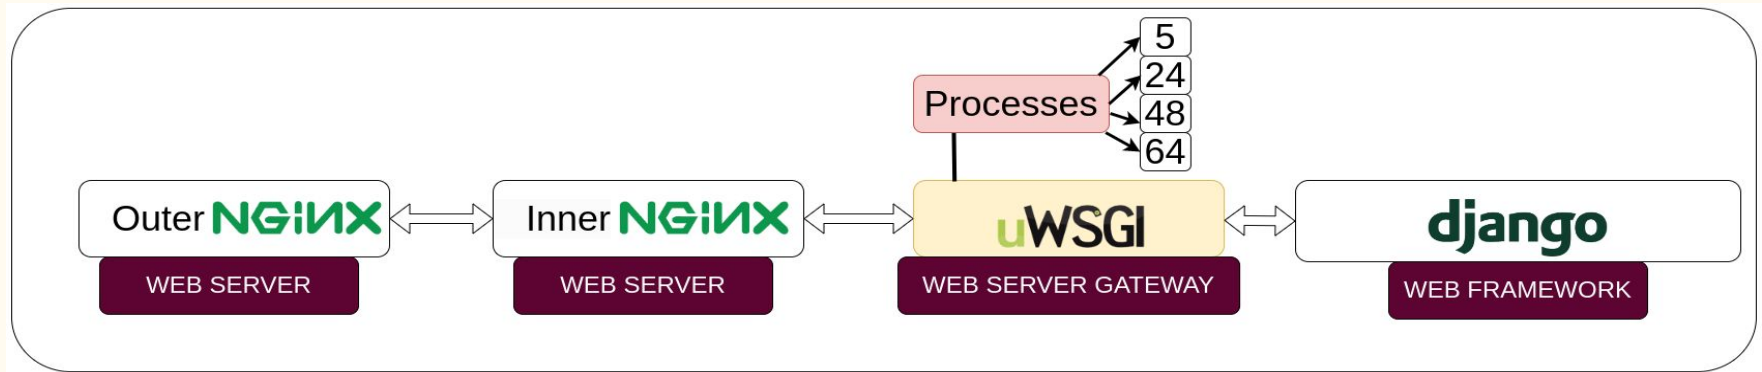
Parameter	Value
RAM	128 GB
Architecture	x86_64
No of CPU	24
Model name	Intel(R) Xeon(R) CPU E5-2650 v4 @ 2.20GHz
Cores per socket	12
Thread per core	2

Machine : safev2 is hosted

# SPC Usage Testing

## Test Setup | uWSGI configurations

We tried **four configurations** for number of processes in **uWSGI**




# SPC Usage Testing

## Test Setup | uWSGI configurations

5: lower than ideal config. | 24: equal to num of cores

48: ideal config. (documentation) | 64: higher than ideal config.

Test: `$ ./client_end_script -l 50 -u 400 -s 50 -t 60`






# SPC Usage Testing


## Final Test | uWSGI Processes 5

---



# SPC Usage Testing

## Final Test | uWSGI Processes 5




This wide gap indicates a bottleneck.

# SPC Usage Testing

## Final Test | uWSGI Processes 24


---



# SPC Usage Testing

## Final Test | uWSGI Processes 48


---



# SPC Usage Testing

## Final Test | uWSGI Processes 64

---



# SPC Usage Testing

## Final Test | Results

---

RESULTS	Worst Performance	Best Performance
Low load(50 users)	uWSGI Processes : 64 Response time : 141.7 ms	uWSGI Processes : 24 Response time : 26.9 ms
High load(400 users)	uWSGI Processes : 5 Response time : 402 ms	uWSGI Processes : 48 Response time : 27 ms

# Conclusion & Future Work

## Conclusion

---

SPC works and it helps in bottleneck detection

It satisfies the following conditions :

independent of any application specific APIs

easy to use and configure

visual feedback for analysis of components.



# Conclusion & Future Work

## Conclusion | Other use cases

---

To fine tune application performance

To check impact of certain API endpoints by modifying  
/sys\_perf\_check API

# Conclusion & Future Work

## Future Work

---

Figure out ways to **extend SPC for more use cases.**

Experimenting with **different implementations of sys\_perf\_check Endpoint.**

More **user friendly naming for test-id** names instead of metrohash

Extend it **for non Linear stacks** as well.

The End